

hw_2_3_v2

Grigoreva Elizaveta

6/7/2020

```
library(mlbench)
library(ggplot2)
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(caret)

## Loading required package: lattice

library(boot)

##
## Attaching package: 'boot'

## The following object is masked from 'package:lattice':
##
##   melanoma

library(tree)
library(class)
library(corrplot)

## corrplot 0.84 loaded

library(ggfortify)
library(RColorBrewer)
library(data.table)

##
## Attaching package: 'data.table'

## The following objects are masked from 'package:dplyr':
##
##   between, first, last

Load data, check NA, duplicates
```

```
data("BostonHousing")
boston <- BostonHousing
str(boston)
```


```
## 'data.frame': 506 obs. of 14 variables:
## $ crim : num 0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn : num 18 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus : num 2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas : Factor w/ 2 levels "0","1": 1 1 1 1 1 1 1 1 1 1 ...
## $ nox : num 0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm : num 6.58 6.42 7.18 7 7.15 ...
## $ age : num 65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis : num 4.09 4.97 4.97 6.06 6.06 ...
## $ rad : num 1 2 2 3 3 3 5 5 5 ...
## $ tax : num 296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num 15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ b : num 397 397 393 395 397 ...
## $ lstat : num 4.98 9.14 4.03 2.94 5.33 ...
## $ medv : num 24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
summary(boston)
```

```
##      crim          zn          indus      chas          nox
## Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   0:471   Min.   :0.3850
## 1st Qu.: 0.08204   1st Qu.: 0.00   1st Qu.: 5.19   1: 35   1st Qu.:0.4490
## Median : 0.25651   Median : 0.00   Median : 9.69           Median :0.5380
## Mean   : 3.61352   Mean   : 11.36   Mean   :11.14           Mean   :0.5547
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10           3rd Qu.:0.6240
## Max.   :88.97620   Max.   :100.00   Max.   :27.74           Max.   :0.8710
##      rm          age          dis          rad
## Min.   :3.561   Min.   : 2.90   Min.   : 1.130   Min.   : 1.000
## 1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100   1st Qu.: 4.000
## Median :6.208   Median : 77.50   Median : 3.207   Median : 5.000
## Mean   :6.285   Mean   : 68.57   Mean   : 3.795   Mean   : 9.549
## 3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188   3rd Qu.:24.000
## Max.   :8.780   Max.   :100.00   Max.   :12.127   Max.   :24.000
##      tax          ptratio          b          lstat
## Min.   :187.0   Min.   :12.60   Min.   : 0.32   Min.   : 1.73
## 1st Qu.:279.0   1st Qu.:17.40   1st Qu.:375.38   1st Qu.: 6.95
## Median :330.0   Median :19.05   Median :391.44   Median :11.36
## Mean   :408.2   Mean   :18.46   Mean   :356.67   Mean   :12.65
## 3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:396.23   3rd Qu.:16.95
## Max.   :711.0   Max.   :22.00   Max.   :396.90   Max.   :37.97
##      medv
## Min.   : 5.00
## 1st Qu.:17.02
## Median :21.20
## Mean   :22.53
## 3rd Qu.:25.00
## Max.   :50.00
```

```
hist(boston$medv)
```

Histogram of boston\$medv



```
summary(is.na(boston))
```

```
##      crim      zn      indus      chas
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506     FALSE:506     FALSE:506     FALSE:506
##      nox      rm      age      dis
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506     FALSE:506     FALSE:506     FALSE:506
##      rad      tax      ptratio      b
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506     FALSE:506     FALSE:506     FALSE:506
##      lstat      medv
## Mode :logical  Mode :logical
## FALSE:506     FALSE:506
```

```
summary(duplicated(boston))
```

```
##      Mode      FALSE
## logical      506
```

Create subset w/o Charles River dummy variable and perform correlation plot

```
summary(is.na(boston))
```

```
##      crim      zn      indus      chas
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506     FALSE:506     FALSE:506     FALSE:506
##      nox      rm      age      dis
## Mode :logical Mode :logical Mode :logical Mode :logical
## FALSE:506     FALSE:506     FALSE:506     FALSE:506
##      rad      tax      ptratio      b
## Mode :logical Mode :logical Mode :logical Mode :logical
```

```
## FALSE:506      FALSE:506      FALSE:506      FALSE:506
## lstat          medv
## Mode :logical  Mode :logical
## FALSE:506      FALSE:506
```

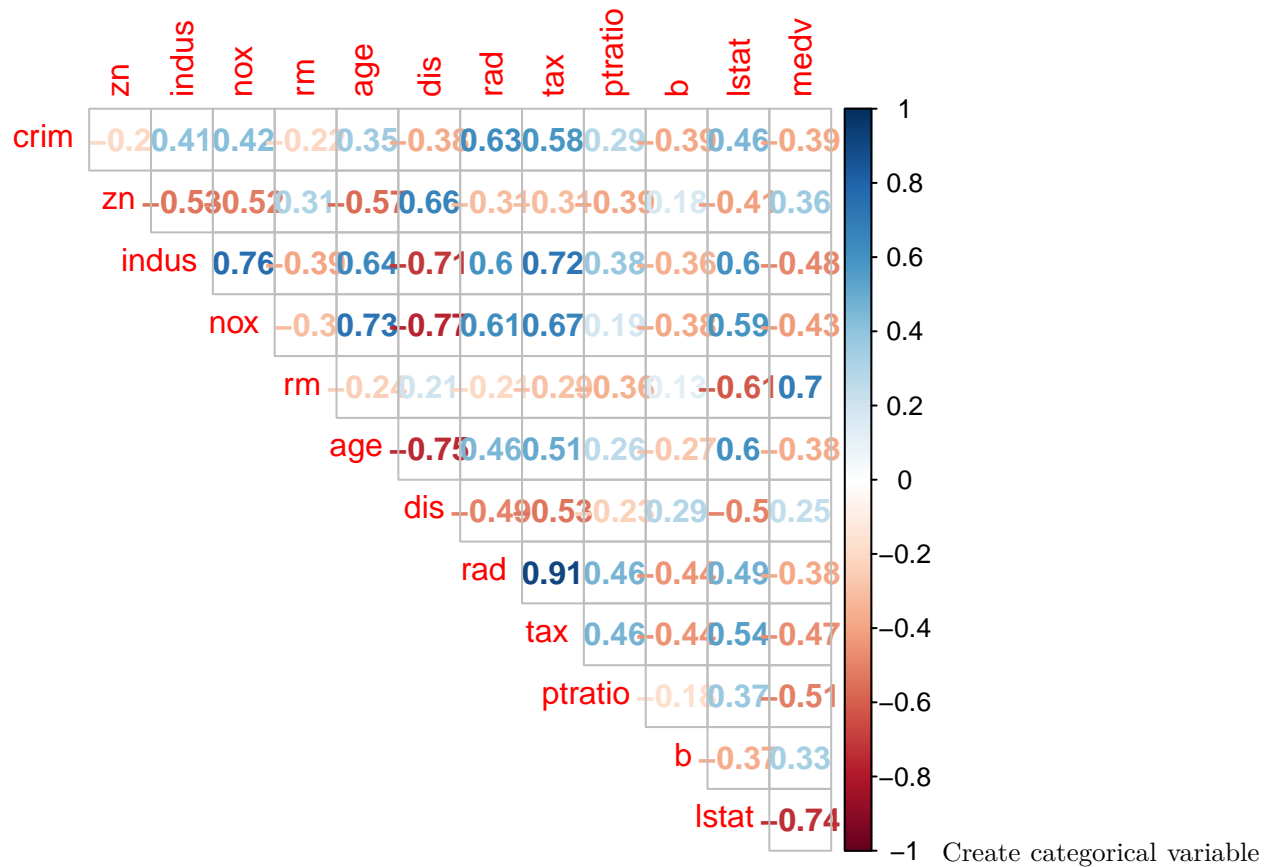
```
summary(duplicated(boston))
```

```
## Mode FALSE
## logical 506
```

```
#Remove Charles River variable
```

```
boston_sub<- boston[,-4]
```

```
corrplot(cor(boston_sub), method="number", type="upper", diag=F)
```



```
randomly
```

```
boston_sub1 <- boston_sub
```

```
boston_sub1 <- data.table(boston_sub1)
```

```
boston_sub1[,medv_Cat:= cut(medv,c(0,quantile(boston_sub1$medv, 0.25),quantile(boston_sub1$medv, 0.50),
```

```
#Table it
```

```
boston_sub1[,table(medv_Cat)]
```

```
## medv_Cat
```

```
## 1 2 3 4
```

```
## 127 129 126 124
```


```
summary(boston_sub1)
```

```
##      crim          zn          indus          nox
## Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   Min.   :0.3850
## 1st Qu.: 0.08204   1st Qu.: 0.00   1st Qu.: 5.19   1st Qu.:0.4490
## Median : 0.25651   Median : 0.00   Median : 9.69   Median :0.5380
## Mean   : 3.61352   Mean    : 11.36   Mean    :11.14   Mean    :0.5547
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10   3rd Qu.:0.6240
## Max.   :88.97620   Max.    :100.00   Max.    :27.74   Max.    :0.8710
##      rm          age          dis          rad
## Min.   :3.561     Min.   : 2.90   Min.   : 1.130   Min.   : 1.000
## 1st Qu.:5.886     1st Qu.: 45.02   1st Qu.: 2.100   1st Qu.: 4.000
## Median :6.208     Median : 77.50   Median : 3.207   Median : 5.000
## Mean   :6.285     Mean    : 68.57   Mean    : 3.795   Mean    : 9.549
## 3rd Qu.:6.623     3rd Qu.: 94.08   3rd Qu.: 5.188   3rd Qu.:24.000
## Max.   :8.780     Max.    :100.00   Max.    :12.127   Max.    :24.000
##      tax          ptratio          b          lstat
## Min.   :187.0     Min.   :12.60   Min.   : 0.32   Min.   : 1.73
## 1st Qu.:279.0     1st Qu.:17.40   1st Qu.:375.38   1st Qu.: 6.95
## Median :330.0     Median :19.05   Median :391.44   Median :11.36
## Mean   :408.2     Mean    :18.46   Mean    :356.67   Mean    :12.65
## 3rd Qu.:666.0     3rd Qu.:20.20   3rd Qu.:396.23   3rd Qu.:16.95
## Max.   :711.0     Max.    :22.00   Max.    :396.90   Max.    :37.97
##      medv      medv_Cat
## Min.   : 5.00     1:127
## 1st Qu.:17.02     2:129
## Median :21.20     3:126
## Mean   :22.53     4:124
## 3rd Qu.:25.00
## Max.   :50.00
```

```
boston_df1 <- boston_sub1[, -13]
#Perform PCA based on this classification
```

PCA

```
boston_pca <- boston_df1[, -13]
pca <- prcomp(boston_pca, center = TRUE,
              scale. = TRUE)
autoplot(pca, data = boston_df1, colour = "medv_Cat")
```



```
#So we have 5 categories but 2 clear clusters, let's make 2 clusters
```

```
#Try to set up 2 clusters
```

```
boston_sub <- data.table(boston_sub)
```

```
boston_sub[,medv_Cat:= cut(medv,c(0,quantile(boston_sub$medv, 0.35), max(boston_sub$medv)),labels=c("yes", "no"))]
```

```
boston_sub[,table(medv_Cat)]
```

```
## medv_Cat
```

```
## yes no
```


```
## 179 327
```

```
boston_df <- boston_sub[, -13]
```

```
boston_pca <- boston_df[, -13]
```

```
pca <- prcomp(boston_pca, center = TRUE,  
              scale. = TRUE)
```

```
autoplot(pca, data = boston_df, colour = "medv_Cat")
```



KNN for predict to which category some observation belongs to

```
set.seed(42)
```

```
for_train <- createDataPartition(y = boston_sub$medv_Cat, p= 0.8, list = FALSE)
training <- boston_sub[for_train,-13]
testing <- boston_sub[-for_train,-13]
```

```
knn <- train(medv_Cat ~ ., data = training, method = "knn",
  trControl=trainControl(method = "repeatedcv", number = 10, repeats = 3),
  preProcess = c("center", "scale"),
  tuneLength = 20) #Range of k?
```

```
knn
```

```
## k-Nearest Neighbors
```

```
##
```

```
## 406 samples
```

```
## 12 predictor
```

```
## 2 classes: 'yes', 'no'
```

```
##
```

```
## Pre-processing: centered (12), scaled (12)
```

```
## Resampling: Cross-Validated (10 fold, repeated 3 times)
```


```
## Summary of sample sizes: 365, 366, 366, 365, 365, 365, ...
```

```
## Resampling results across tuning parameters:
```

```
##
```

```
## k Accuracy Kappa
## 5 0.8767480 0.7241543
## 7 0.8693699 0.7097523
## 9 0.8652236 0.7025707
## 11 0.8570528 0.6839796
## 13 0.8456301 0.6573275
## 15 0.8430488 0.6521292
## 17 0.8323577 0.6273424
## 19 0.8283130 0.6198062
## 21 0.8341667 0.6307289
## 23 0.8324797 0.6239412
## 25 0.8325407 0.6227029
## 27 0.8358333 0.6314846
## 29 0.8391870 0.6382910
## 31 0.8333740 0.6251189
## 33 0.8325407 0.6235014
## 35 0.8350000 0.6295236
## 37 0.8317480 0.6216412
## 39 0.8227236 0.6018037
## 41 0.8186179 0.5934927
## 43 0.8096341 0.5736192
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 5.
```

```
plot(knn)
```



```
print(paste0("best k to minimize MSE: ", knn$bestTune))
```

```
## [1] "best k to minimize MSE: 5"
```



```

#Test this model
test_pred <- predict(knn, newdata =testing)
table(test_pred, Real = testing$medv_Cat)

##           Real
## test_pred yes no
##      yes  30  3
##      no   5 62

knn_fit <- knn3Train(train = training[,-13], test = testing[,-13], k=7, cl = training$medv_Cat)
table(knn_fit, Real = testing$medv_Cat)

##           Real
## knn_fit yes no
##      no   9 59
##      yes  26 6

#accuracy is a ratio of correctly predicted observation to the total observations. (TP+TN)/(FP+FN+TP+TN)
accuracy = sum(knn_fit == testing$medv_Cat)/length(testing$medv_Cat)
paste0("Accuracy:", accuracy)

## [1] "Accuracy:0.85"

Logistic regression

#Check multicol
mean(boston$medv)


## [1] 22.53281

#Create binomial variable and separate to the hight and low level of criminal based on median

ggplot(boston,aes(lstat,medv))+geom_point()+geom_smooth()

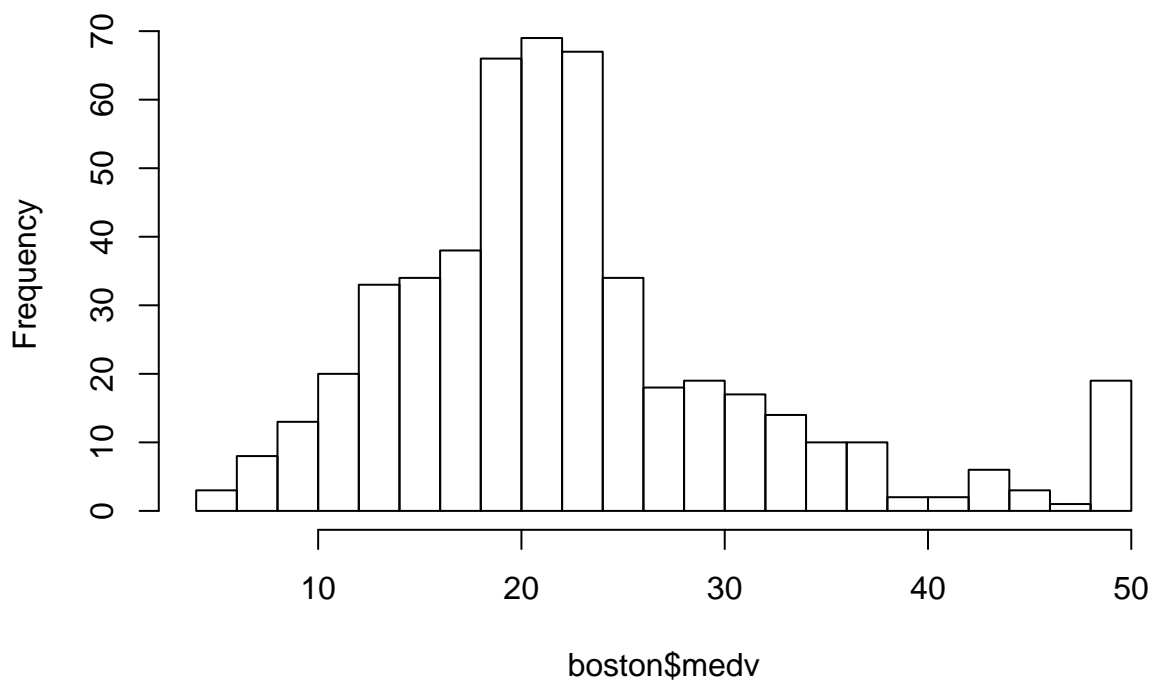
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'

```



```
cv.err <- 1:5
hist(boston$medv, median(boston$medv))
```

Histogram of boston\$medv




```
median(boston$medv)
```

```
## [1] 21.2
```

```
boston_bi <- boston %>% mutate(medv_bi = ifelse(medv >= 22.5, "high", "low")) %>%  
  mutate(medv_bi = factor(medv_bi, levels = c("high", "low")))  
boston_bi <- boston_bi[, -c(13:14)]  
for (i in 1:5){  
  gl <- glm(medv_bi ~ poly(rm, i), family = "binomial",  
            data = boston_bi)  
  cv.err[i] <- cv.glm(boston_bi, gl)$delta[1]  
}  
cv.err
```

```
## [1] 0.1576716 0.1563841 0.1529078 0.1534497 0.1537524
```


```
plot(x = 1:5, y = cv.err,  
     xlab = 'Polynomial degree', ylab='Cross-validation')
```



```
# on each iteration, create the model for the given power and test MSE  
for (i in 1:5){  
  gl <- glm(medv_bi ~ poly(rm, i), family = "binomial",  
            data = boston_bi)  
  cv.err[i] <- cv.glm(boston_bi, gl)$delta[1]  
}  
cv.err
```

```
## [1] 0.1576716 0.1563841 0.1529078 0.1534497 0.1537524
```

```
plot(x = 1:5, y = cv.err,  
     xlab = 'Polynomial degree', ylab='Cross-validation')
```



```
# use 3rd degree
glm <- glm(medv_bi ~ poly(rm, 3), data = boston_bi, family = "binomial")
summary(glm)

##
## Call:
## glm(formula = medv_bi ~ poly(rm, 3), family = "binomial", data = boston_bi)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.5703  -0.6721   0.3852   0.7137   2.4747
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    0.3929     0.1198   3.279  0.00104 **
## poly(rm, 3)1 -34.8454     3.5521 -9.810 < 2e-16 ***
## poly(rm, 3)2  -5.0733     3.1524 -1.609  0.10754
## poly(rm, 3)3  18.5221     3.0243   6.124  9.1e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 688.12  on 505  degrees of freedom
## Residual deviance: 469.84  on 502  degrees of freedom
## AIC: 477.84
##
## Number of Fisher Scoring iterations: 5
pred_glm <- predict(glm, type = "response") > 0.5
table(pred_glm, Real = boston_bi[, 13])

##      Real
```

```
## pred_glm high low
## FALSE 146 41
## TRUE 66 253

boston_df <- boston_sub[, -13]
test <- sample(nrow(boston_df), 0.8*nrow(boston_df))
glm2 <- glm(medv_bi ~ rm, data = boston_bi[-test, ],
            family = "binomial")

#Summary
summary(glm2)

##
## Call:
## glm(formula = medv_bi ~ rm, family = "binomial", data = boston_bi[-test,
## ])
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2001  -1.0995   0.6616   0.8873   2.4534
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   9.2227     2.6394   3.494 0.000475 ***
## rm          -1.3874     0.4201  -3.303 0.000958 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 133.62  on 101  degrees of freedom
## Residual deviance: 118.13  on 100  degrees of freedom
## AIC: 122.13
##
## Number of Fisher Scoring iterations: 4

pred_glm2 <- predict(glm2, type = "response", newdata = boston_bi[test, ]) > 0.5
table(pred_glm2, Real = boston_bi[test, 13])

##           Real
## pred_glm2 high low
## FALSE    88  11
## TRUE     87 218
```