

Homework 2.3

Valeriia

20 05 2020

Load all needed libraries.

```
library(MASS)
library(data.table)
library(ggplot2)
library(caret)
library(boot)
library(corrplot)
data(Boston)
```

Check our data. It's ok.

```
bos <- Boston
str(bos)
```

```
## 'data.frame':  506 obs. of  14 variables:
## $ crim   : num  0.00632 0.02731 0.02729 0.03237 0.06905 ...
## $ zn     : num  18 0 0 0 0 0 12.5 12.5 12.5 12.5 ...
## $ indus  : num  2.31 7.07 7.07 2.18 2.18 2.18 7.87 7.87 7.87 7.87 ...
## $ chas   : int   0 0 0 0 0 0 0 0 0 0 ...
## $ nox    : num  0.538 0.469 0.469 0.458 0.458 0.458 0.524 0.524 0.524 0.524 ...
## $ rm     : num  6.58 6.42 7.18 7 7.15 ...
## $ age    : num  65.2 78.9 61.1 45.8 54.2 58.7 66.6 96.1 100 85.9 ...
## $ dis    : num  4.09 4.97 4.97 6.06 6.06 ...
## $ rad    : int   1 2 2 3 3 3 5 5 5 5 ...
## $ tax    : num  296 242 242 222 222 222 311 311 311 311 ...
## $ ptratio: num  15.3 17.8 17.8 18.7 18.7 18.7 15.2 15.2 15.2 15.2 ...
## $ black  : num  397 397 393 395 397 ...
## $ lstat  : num  4.98 9.14 4.03 2.94 5.33 ...
## $ medv   : num  24 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 ...
```

```
dim(bos)
```

```
## [1] 506  14
```

```
sum(is.na(bos))
```

```
## [1] 0
```

```
summary(bos)
```

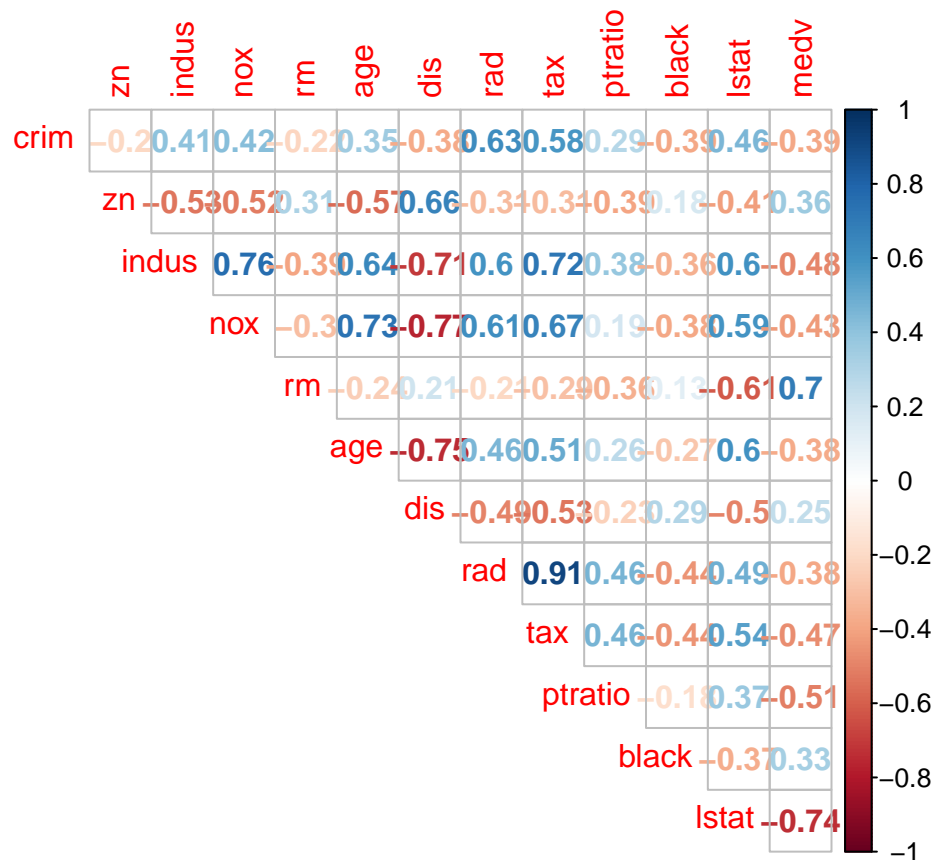
```
##          crim              zn          indus          chas
## Min.   : 0.00632   Min.   : 0.00   Min.   : 0.46   Min.   :0.00000
## 1st Qu.: 0.08204   1st Qu.: 0.00   1st Qu.: 5.19   1st Qu.:0.00000
## Median : 0.25651   Median : 0.00   Median : 9.69   Median :0.00000
## Mean   : 3.61352   Mean   : 11.36   Mean   :11.14   Mean   :0.06917
## 3rd Qu.: 3.67708   3rd Qu.: 12.50   3rd Qu.:18.10   3rd Qu.:0.00000
## Max.   :88.97620   Max.   :100.00   Max.   :27.74   Max.   :1.00000
##          nox          rm          age          dis
## Min.   :0.3850   Min.   :3.561   Min.   : 2.90   Min.   : 1.130
## 1st Qu.:0.4490   1st Qu.:5.886   1st Qu.: 45.02   1st Qu.: 2.100
## Median :0.5380   Median :6.208   Median : 77.50   Median : 3.207
## Mean   :0.5547   Mean   :6.285   Mean   : 68.57   Mean   : 3.795
## 3rd Qu.:0.6240   3rd Qu.:6.623   3rd Qu.: 94.08   3rd Qu.: 5.188
## Max.   :0.8710   Max.   :8.780   Max.   :100.00   Max.   :12.127
##          rad          tax          ptratio          black
## Min.   : 1.000   Min.   :187.0   Min.   :12.60   Min.   : 0.32
## 1st Qu.: 4.000   1st Qu.:279.0   1st Qu.:17.40   1st Qu.:375.38
## Median : 5.000   Median :330.0   Median :19.05   Median :391.44
## Mean   : 9.549   Mean   :408.2   Mean   :18.46   Mean   :356.67
## 3rd Qu.:24.000   3rd Qu.:666.0   3rd Qu.:20.20   3rd Qu.:396.23
## Max.   :24.000   Max.   :711.0   Max.   :22.00   Max.   :396.90
##          lstat          medv
## Min.   : 1.73   Min.   : 5.00
## 1st Qu.: 6.95   1st Qu.:17.02
## Median :11.36   Median :21.20
## Mean   :12.65   Mean   :22.53
## 3rd Qu.:16.95   3rd Qu.:25.00
## Max.   :37.97   Max.   :50.00
```

```
sum(duplicated(bos))
```

```
## [1] 0
```

Charles River dummy variable (= 1 if tract bounds river; 0 otherwise). Drop it. Check correlation.

```
bos <- bos[,-4]
corrplot(cor(bos), method = "number", type = "upper", diag = FALSE)
```



Cathecorize data.

```
bos <- data.table(bos)
bos[,medv_F:= cut(medv,c(0,quantile(bos$medv, 0.33),quantile(bos$medv, 0.66),
max(bos$medv)),labels=c("1","2","3"))]
bos[,table(medv_F)]
```

```
## medv_F
## 1 2 3
## 167 167 172
```


PCA.

```
pca <- prcomp(bos[,c(13,14)], center = TRUE, scale. = TRUE)
pca
```

```
## Standard deviations (1, ..., p=12):
## [1] 2.4752210 1.1586541 1.0861790 0.9138194 0.8152738 0.7330805 0.6296169
## [8] 0.5263720 0.4693245 0.4314643 0.4114793 0.2542551
##
## Rotation (n x k) = (12 x 12):
##          PC1      PC2      PC3      PC4      PC5      PC6
## crim    0.2510194 -0.40124935 0.06905890 -0.07537120 0.20811692 -0.77883466
## zn      -0.2562763 -0.43910123 0.09079765 -0.30453280 0.35173539 0.27063296
## indus    0.3466252 0.10826541 0.03147503 0.01032982 0.09094655 0.34042847
```

```
## nox      0.3427732  0.16538852  0.23634943 -0.14816181  0.14361054  0.19010878
## rm      -0.1893443 -0.07676264  0.67862805  0.39317754 -0.10468673 -0.07757913
## age      0.3135926  0.31549763  0.16446704 -0.03415323  0.04226173 -0.12864962
## dis     -0.3214520 -0.32731960 -0.25420475 -0.07645279  0.01083277  0.11493067
## rad      0.3198193 -0.38437642  0.11313425  0.21734697  0.16394203  0.14004949
## tax      0.3385180 -0.32057097  0.07803058  0.14109504  0.21000125  0.31042396
## ptratio  0.2050739 -0.17273359 -0.48516673  0.60814980 -0.24573826  0.01417231
## black   -0.2030293  0.33625146 -0.18803264  0.36330273  0.81001861 -0.09211058
## lstat    0.3098245  0.03364522 -0.29715025 -0.38592449  0.06096419 -0.08775570
##          PC7          PC8          PC9          PC10          PC11
## crim     0.158230415 -0.26179898  0.01980470 -0.11039262 -0.08663749
## zn       -0.403359872 -0.35858760  0.26689361  0.26335598  0.07080081
## indus    0.173213403 -0.64380852 -0.36378621 -0.30263190  0.11400759
## nox      0.076735083  0.01984964  0.23045723  0.11122392 -0.80409242
## rm      -0.329939431 -0.04637159 -0.43204790  0.05329734 -0.15277946
## age     -0.602218414  0.06657278  0.36352844 -0.45777362  0.21328031
## dis     -0.118900334  0.15077227 -0.16897171 -0.69863534 -0.38969478
## rad      0.080571835  0.46980760  0.02306152  0.03388316  0.10601899
## tax      0.079371592  0.17846255 -0.03592292 -0.10173769  0.21630840
## ptratio -0.313349725 -0.25684385  0.15388089  0.17261849 -0.21044696
## black   -0.008231654  0.04631318 -0.09701329  0.02002712 -0.04151870
## lstat   -0.423864339  0.19492701 -0.60053449  0.26891408 -0.05644192
##          PC12
## crim    -0.044517237
## zn       0.081746955
## indus    0.247896886
## nox     -0.047399129
## rm      -0.047915403
## age      0.035654465
## dis      0.019094950
## rad      0.635066577
## tax     -0.720833262
## ptratio -0.019118760
## black    0.002297948
## lstat   -0.020576836
```

```
ntp <- data.frame('medv_F' = bos$medv_F, pca$x[,1:2])
ggplot(data = ntp) +
  geom_point(aes(x = PC1, y = PC2, col = medv_F)) +
  theme_minimal()
```



KNN.


```
set.seed(3)
intrain <- createDataPartition(y = bos$medv_F, p= 0.8, list = FALSE)
training <- bos[intrain,-13]
testing <- bos[-intrain,-13]
# Repeat several times, 10 blocks, 3 times
trctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 3)

knn_fit <- train(medv_F ~ ., data = training, method = "knn",
  trControl=trctrl,
  preProcess = c("center", "scale"),
  tuneLength = 20)
knn_fit
```

```
## k-Nearest Neighbors
##
## 406 samples
## 12 predictor
## 3 classes: '1', '2', '3'
##
## Pre-processing: centered (12), scaled (12)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 366, 365, 365, 367, 365, ...
## Resampling results across tuning parameters:
##
```

```
## k Accuracy Kappa
## 5 0.7624363 0.6437411
## 7 0.7676367 0.6516739
## 9 0.7668867 0.6506242
## 11 0.7619864 0.6432025
## 13 0.7437962 0.6158727
## 15 0.7461732 0.6193821
## 17 0.7421072 0.6134217
## 19 0.7412952 0.6121960
## 21 0.7330411 0.5998296
## 23 0.7314171 0.5974494
## 25 0.7338958 0.6011902
## 27 0.7324142 0.5989529
## 29 0.7332474 0.6002403
## 31 0.7291398 0.5941370
## 33 0.7209254 0.5818946
## 35 0.7241755 0.5867250
## 37 0.7192171 0.5792872
## 39 0.7134030 0.5706491
## 41 0.7059630 0.5595450
## 43 0.7051114 0.5583733
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was k = 7.
```

```
# The hieghest accuracy is in k=7
plot(knn_fit)
```



```
test_pred <- predict(knn_fit, newdata =testing)
head(data.frame(test_pred, testing$medv_F))
```

```
##   test_pred testing.medv_F
## 1         2             2
## 2         1             2
## 3         2             2
## 4         2             2
## 5         3             2
## 6         2             1
```

```
table(test_pred, Real = testing$medv_F)
```

```
##           Real
## test_pred  1  2  3
##           1 29  6  1
##           2  4 20  7
##           3  0  7 26
```

```
knn_fit <- knn3Train(train = training[,-13], test = testing[,-13], k=7, cl = training$medv_F)
xtab <- table(knn_fit, Real = testing$medv_F)
xtab
```

```
##           Real
```

```
## knn_fit  1  2  3
##          1 29 10  4
##          2  2 15 12
##          3  2  8 18
```

```
accuracy = sum(knn_fit == testing$medv_F)/length(testing$medv_F)
precision = xtab[1,1]/sum(xtab[,1])
recall = xtab[1,1]/sum(xtab[1,])
f = 2 * (precision * recall) / (precision + recall)
```

*#Accuracy - Accuracy is the most intuitive performance measure and it
#is simply a ratio of correctly predicted observation to the total observations. $(TP+TN)/(FP+FN+TP+TN)$*

```
paste0("Accuracy:", accuracy)
```

```
## [1] "Accuracy:0.62"
```

*#Precision - Precision is the ratio of correctly predicted
#positive observations to the total predicted positive observations. $TP/(TP+FP)$*

```
paste0("Precision:", precision)
```

```
## [1] "Precision:0.878787878787879"
```

*#Recall (Sensitivity) - Recall is the ratio of correctly predicted positive
#observations to the all observations in actual class - yes. $TP/(TP+FN)$*

```
paste0("Recall:", recall)
```

```
## [1] "Recall:0.674418604651163"
```


*#F1 Score = $2 * (Recall * Precision) / (Recall + Precision)$*

```
paste0("F:", f)
```

```
## [1] "F:0.763157894736842"
```

Polynomial regression


```
qplot(lstat, medv, data = Boston, geom = c("point", "smooth"), method = "lm")
```

```
bos[,medv_B:= cut(medv,c(0,median(bos$medv),max(bos$medv)),labels=c("low","high"))]
bos[,table(medv_B)]
```

```
## medv_B
## low high
## 256 250
```

```
set.seed(3)
intrain <- createDataPartition(y = bos$medv_B, p= 0.8, list = FALSE)
bos1 <- bos[,-c(13,14)]
training <- bos[intrain,-c(13,14)]
testing <- bos[-intrain,-c(13,14)]
#k-fold CV
errors <- c()
for (i in 1:5){
  g <- glm(medv_B ~ poly(lstat,i), family = "binomial", data = bos1)
  errors[i] <- cv.glm(bos1, g)$delta[1]
}
plot(x = 1:5, y = errors, xlab = 'Polynomial degree', ylab = 'Cross-validation')
```



```
mod <- glm(medv_B ~ poly(lstat, 5), data = training, family = "binomial")
summary(mod)
```

```
##
## Call:
## glm(formula = medv_B ~ poly(lstat, 5), family = "binomial", data = training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.0865  -0.5115   0.0000   0.5056   2.2062
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -16.29     13.27  -1.227   0.220
## poly(lstat, 5)1 -1131.94    832.37  -1.360   0.174
## poly(lstat, 5)2 -1394.55   1094.94  -1.274   0.203
## poly(lstat, 5)3 -1241.24    925.87  -1.341   0.180
## poly(lstat, 5)4  -613.34    461.16  -1.330   0.184
## poly(lstat, 5)5  -221.79    154.13  -1.439   0.150
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 561.39  on 404  degrees of freedom
## Residual deviance: 284.08  on 399  degrees of freedom
## AIC: 296.08
```

```
##
## Number of Fisher Scoring iterations: 15
```

```
pred <- predict(mod, type = 'response') > 0.5
table(pred, Real = training$medv_B)
```

```
##          Real
## pred    low high
##  FALSE  170   33
##   TRUE   35  167
```

```
mod2 <- glm(medv_B ~ lstat, data = testing, family = "binomial")
summary(mod2)
```

```
##
## Call:
## glm(formula = medv_B ~ lstat, family = "binomial", data = testing)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6348  -0.5948  -0.0301   0.5975   3.2198
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  3.66041    0.73695   4.967 6.80e-07 ***
## lstat       -0.29910    0.06008  -4.978 6.41e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 140.006  on 100  degrees of freedom
## Residual deviance:  86.562  on  99  degrees of freedom
## AIC: 90.562
##
## Number of Fisher Scoring iterations: 5
```

```
pred2 <- predict(mod2, type = 'response', newdata = testing) > 0.5
table(pred2, Real = testing$medv_B)
```

```
##          Real
## pred2    low high
##  FALSE   39    8
##   TRUE   12   42
```